

Stakeholder Engagement Plan

Latent Probing for Mental Health Sentiment Classification

This plan identifies two stakeholder groups relevant to our project, the potential benefits and risks our work presents to each, the engagement method selected for each group from the Action Catalogue, and how feedback from these groups could influence our dataset, model design, outputs, evaluation, and deployment decisions.

1. Stakeholder Group: Mental Health Practitioners / Counselors

Who they are

Licensed clinicians, therapists, and counselors, including the psychology or counseling professional our team plans to consult to review our DSM-5-based labeling rubric, as described in our Data Card and Impact Statement.

Why they are important to our project

They are the domain experts best positioned to judge whether our Depression, Anxiety, Stress, and Normal labels and rubric reflect clinically sound indicators, rather than only statistical patterns learned from social media text.

What matters most to them

- Clinical accuracy and avoiding false reassurance or false alarm.
- The tool being positioned as a support aid, not a diagnostic replacement, in line with our stated safety value.
- Their professional judgment remaining central to any decision made using the model's output.

2. Potential Benefits and Risks

Potential benefits

- Improves the rigor and credibility of our labeling process.
- Gives practitioners a lightweight screening aid that could help school counseling programs or peer support groups direct limited attention toward people who may need it most.

Potential harms or concerns

- Risk of the tool being over-trusted as if it were a diagnosis.
- Risk of practitioner time being spent validating a rubric for a tool not yet ready for real-world deployment.
- Reputational risk to practitioners if associated with a tool that underperforms for certain populations.

Ethical, social, cultural, or economic considerations

Practitioner involvement supports our safety value by keeping human judgment central to any decision made using model output, and helps guard against the model being treated as an autonomous diagnostic authority.

3. Engagement Method

Method selected: User Committees

Why the method is appropriate: A small, recurring review committee gives ongoing, structured input rather than a single one-off consultation, which suits a rubric that may need revisiting as labeling or model choices evolve.

Why it is suitable for this group: It mirrors the consultation step our team already committed to in our Data Card and Impact Statement, and matches how practitioners are typically engaged, as advisors reviewing technical work rather than co-designers of it.

How the engagement would be conducted: One to two mental health professionals would be invited to review our labeling rubric and a sample of dataset labels, flag any labels that do not match clinical indicators, and provide feedback at key project checkpoints.

2. Stakeholder Group: Underrepresented Community Members

Who they are

Non-English speakers, older adults, people without regular internet or social media access, and individuals who avoid posting publicly about mental health due to stigma, the groups our own Data Card identifies as likely underrepresented in our training data.

Why they are important to our project

They represent the populations most at risk of the tool performing poorly or unfairly for them, which directly bears on our stated fairness value. Our Data Card notes that silent sufferers may be systematically excluded from the very dataset meant to help detect them.

What matters most to them

- Not being further marginalized by a tool that does not recognize their language, context, or way of expressing distress.
- Having their absence from the dataset acknowledged transparently rather than hidden.
- Meaningful engagement rather than token consultation.

2. Potential Benefits and Risks

Potential benefits

If engaged, their input could guide future work on expanding language and demographic coverage, a mitigation step already recommended in our Data Card.

Potential harms or concerns

- If not engaged, the tool could give false confidence in how broadly it generalizes.
- Risk of tokenizing this group if engagement is shallow rather than genuine and ongoing.

Ethical, social, cultural, or economic considerations

Engaging this group supports our fairness value by ensuring known representation gaps are documented openly rather than overstated as generalizable results, consistent with our Data Card's

bias and representation section.

3. Engagement Method

Method selected: Distributed Dialogue

Why the method is appropriate: Because this group is, by definition, hard to reach through our existing dataset or a single centralized meeting, an asynchronous, distributed approach is more realistic than a single workshop.

Why it is suitable for this group: It lowers the barrier to participation for people with limited internet access, language differences, or reluctance to speak about mental health in a group setting.

How the engagement would be conducted: Short, asynchronous feedback channels, such as simple surveys, shared through community organizations, student groups, or organizations working with older adults or non-English speakers, inviting reflection on gaps in the dataset and tool.

4. Influence on Project Design

Feedback from these two stakeholder groups is expected to influence our project as follows:

- **Dataset:** Practitioner feedback helps validate label quality; community feedback informs what future data collection should prioritize, including languages and demographics.
- **Model design:** Practitioner input on which distress indicators matter clinically could help refine which latent layers or features our probe focuses on.
- **Outputs:** Both groups reinforce our decision to output confidence or probability distributions rather than a single hard label, since a soft, interpretable signal is safer to act on than a false-certain diagnosis.
- **Evaluation:** Practitioner-validated relabeling of a sample of posts, using our DSM-5-based rubric, becomes part of our evaluation process, rather than relying only on the original dataset labels.
- **Deployment decisions:** Concerns raised by both stakeholder groups support our decision to keep this project at research and educational scope only, rather than pursuing public deployment at this stage, consistent with our Impact Statement.

This plan was prepared as part of the Cohort 10 project submission requirements and reflects commitments already documented in our Problem Statement, Data Card, and Impact Statement.